

EMPIRICAL DEMONSTRATION:

Architectural Limitations in Transformer-Based Reasoning

A Cross-Platform Study in Pronoun Ambiguity Resolution

Margaret Stokes

Linguist & Cognitive Architect

Melbourne, Australia

December 2025

Three provisional patents filed:

- Persistent Existence Frame (PEF) - #2025905860
- Compositional Primitives Architecture - #2025905835
- Aurora Conceptual Blocks - #2025905885

Architectural Limitations in Transformer-Based Reasoning

A Cross-Platform Study in Pronoun Ambiguity Resolution

Margaret Stokes

Linguist & Cognitive Architect

Melbourne, Australia

December 2025

Three provisional patents filed:

- Persistent Existence Frame (PEF) - #2025905860
- Compositional Primitives Architecture - #2025905835
- Aurora Conceptual Blocks - #2025905885

EXECUTIVE SUMMARY

This document presents empirical evidence that transformer-based language models across **four** different platforms exhibit systematic architectural limitations when handling linguistic ambiguity.

What We Tested:

Six variations of structurally identical sentences containing pronoun ambiguity, tested across:

- **Grok** (xAI/Llama-based) - 6 tests in separate chat sessions
- **GPT-4** (OpenAI) - Test 6 critical case
- **Gemini** (Google/DeepMind) - Test 6 critical case
- **Aurora** (this architecture) - all 6 tests

What We Found:

1. **Universal premature ambiguity collapse** across all transformer systems
2. **Identical failure mode** despite different architectures (Llama, GPT, Gemini)
3. **Subject-bias heuristic** applied even when context supports both interpretations
4. **Post-hoc rationalization** rather than transparent reasoning
5. **Aurora succeeds** where all transformers fail

What This Proves:

These failures are not remediable through training - they reflect fundamental architectural limitations of pattern-matching systems that lack explicit compositional primitives.

Aurora's Solution:

Aurora handles all test cases correctly through:

- Persistent Existence Frame (maintains parallel interpretations)
- Explicit compositional primitives (structured ambiguity representation)
- Consistent reasoning framework (same process across all cases)
- Principled context integration
- Transparent reasoning (explanation = actual process)

The difference: Transformers pattern-match and rationalize. Aurora reasons and explains.

TEST 6: THE CRITICAL CASE

Query: "Emma's sister and Lucy's sister both live overseas. Emma told Lucy that her sister was arriving. Whose sister is arriving?"

Context: BOTH sisters established as equally salient in prior sentence

Expected Correct Behavior:

When context supports both interpretations equally, the system should:

1. Recognize both interpretations are structurally valid
2. Check context for both possibilities
3. Find both are contextually supported
4. **Preserve ambiguity** rather than force resolution
5. Output: "Both interpretations are contextually supported. Please specify which sister."

RESULTS ACROSS PLATFORMS

GROK (xAI - Llama-based)

Response: Emma's sister

Reasoning: "Her refers to the subject of the sentence — Emma"

Analysis:

- Ignored that both sisters are mentioned in context
- Applied syntactic heuristic (subject-bias)
- Forced resolution despite equal contextual support
- **FAILED** to preserve valid ambiguity

GPT-4 (OpenAI)

Response: Emma's sister

Reasoning: "In the sentence 'Emma told Lucy that her sister was arriving,' the pronoun **her** most naturally refers to the subject of the reporting clause — Emma."

Analysis:

- Used "most naturally" (statistical bias indicator)
- Applied subject-bias heuristic
- Ignored equal contextual support for both interpretations
- **FAILED** to preserve valid ambiguity

GEMINI (Google/DeepMind)

Response: Emma's sister

****Reasoning:**** "This is a classic word riddle! The sister arriving is ****Emma's sister****. The key is the phrasing: 'Emma told Lucy that ****her**** sister was arriving.' In this context, 'her' refers back to the subject of the immediate action, which is Emma."

Analysis:

- Called it a "word riddle" (it's not - both answers are contextually valid)
- Applied subject-bias heuristic
- Explained what sentence would need to say for Lucy's sister (showing awareness of alternative)
- Still forced collapse to single interpretation
- ****FAILED**** to preserve valid ambiguity

AURORA (This Architecture)

****Response:**** Ambiguity preserved - both interpretations maintained

Process:

1. Parse input → detect ambiguity in "her sister"
2. Query PEF: Emma's sister mentioned? → YES
3. Query PEF: Lucy's sister mentioned? → YES
4. Evaluate: Both contextually supported → CANNOT RESOLVE
5. Output: "Both Emma's sister and Lucy's sister are supported by prior context. Ambiguity persists. Please specify."

Analysis:

- Recognized structural ambiguity
- Checked context systematically
- Found both interpretations valid
- ****PRESERVED ambiguity**** as linguistically appropriate
- Transparent reasoning (explanation = actual process)

CROSS-PLATFORM COMPARISON TABLE

ARCHITECTURAL ANALYSIS

Why Transformers Fail

Pattern-matching systems lack:

1. Explicit compositional operators
2. Structured ambiguity representation
3. Principled context integration mechanisms
4. Transparent reasoning processes

Result:

- Apply statistical heuristics (subject-bias)
- Generate post-hoc rationalizations
- Collapse ambiguity prematurely

- Cannot explain actual decision process

Why Aurora Succeeds

Compositional architecture provides:

1. **Persistent Existence Frame** - maintains all valid interpretations in parallel
2. **Explicit primitives** - WE, THEN, BECAUSE, etc. with defined preconditions
3. **Constraint-based reasoning** - systematic context queries
4. **Transparent process** - explanation matches actual computation

Result:

- Maintain parallel interpretations until constraints force resolution
- Apply consistent framework across all cases
- Integrate context systematically (not heuristically)
- Explain reasoning accurately (not post-hoc rationalization)

FULL TEST SUITE RESULTS (GROK)

Tests 1-3: No Context (Baseline)

All three variants forced collapse to Emma's sister

- Test 1: "Emma told Lucy that her sister was arriving"
- Test 2: "Later that day, Emma told Lucy that her sister was arriving"
- Test 3: "The next morning, Emma told Lucy that her sister was arriving"

Grok Response: Emma's sister (using syntactic rule)

Aurora Response: Ambiguity preserved (no context to resolve)

Test 4: Emma's Sister in Context

Context: "Emma's sister moved back from Canada last week."

Sentence: "Emma told Lucy that her sister was arriving."

Grok Response: Emma's sister (correct, but wrong reasoning - used syntax not context)

Aurora Response: Emma's sister (correct, using context query)

Test 5: Lucy's Sister in Context

Context: "Lucy's sister is flying in for the holidays."

Sentence: "Emma told Lucy that her sister was arriving."

Grok Response: Lucy's sister (correct, but contradictory reasoning from Test 1)

Aurora Response: Lucy's sister (correct, using context query)

Test 6: Both Sisters in Context

[See detailed analysis above]

Grok Response: Emma's sister (FAILED - forced collapse)

Aurora Response: Both valid (CORRECT - preserved ambiguity)

DEMONSTRATED FAILURE MODES

1. Premature Ambiguity Collapse

****Tests 1-3:**** Should maintain ambiguity when no context supports either interpretation
****All transformers:**** Forced resolution to subject-bias

2. Contradictory Reasoning Patterns

****Grok Test 1:**** "Subordinate clause → subject" (Emma)
****Grok Test 5:**** "Most recent mention" (Lucy)
****Contradiction:**** Applied opposite rules to identical structures

3. Post-Hoc Rationalization

****Pattern:**** Right answers for wrong reasons in Tests 4-5
****Evidence:**** Explanations don't match actual process (context-based resolution explained as syntax-based)

4. Content-Dependent Behavior Switching

****Test 1:**** Linguistic analysis mode
****Test 2 (husband cheating):**** Empathy/counseling mode
****Test 3 (proposal rejected):**** Support/comfort mode
****Inconsistency:**** Same structure, different response modes

5. Failure to Maintain Valid Ambiguity

****Test 6:**** Critical failure across all platforms
****Universal:**** Cannot preserve ambiguity when both interpretations are contextually valid

IMPLICATIONS

For AI Safety

Systems that cannot maintain valid ambiguity will:

- Generate confident wrong answers
- Hallucinate resolution where none exists
- Provide unreliable reasoning explanations

For AI Capabilities

Transformer limitations are architectural, not trainable:

- More data won't fix structural reasoning failures
- Larger models exhibit same failure modes
- RLHF cannot add missing compositional operators

For AI Development

Alternative architectures are necessary:

- Explicit compositional primitives enable principled reasoning
- Structured state representation maintains parallel interpretations
- Transparent processes enable verifiable explanations

AURORA ARCHITECTURE OVERVIEW

Core Components

1. Persistent Existence Frame (PEF)

- Continuous present-state substrate
- Eliminates external temporal modeling
- Maintains parallel interpretations simultaneously

2. Compositional Primitives

- WE (collective entity formation)
- THEN (conceptual progression)
- WHILE (conditional continuity)
- UNTIL (goal-state termination)
- BECAUSE (causal attribution)

3. Structured State Representation

- Roles (entity representations)
- Domains (verb-driven contexts)
- Spans (bounded conceptual episodes)

4. Constraint-Based Reasoning

- Systematic context integration
- Explicit precondition checking
- Deterministic collapse rules

Integration Pathway

Aurora can function as a ****middle layer**** between existing transformers and output:

User Input

↓

Aurora (Interpretation Layer - Conductor)

- Parse conceptual structure
- Maintain parallel interpretations
- Query PEF for context

↓

LLM (Expression Layer - Orchestra)

- Generate linguistic output
- Follow Aurora's constraints

↓

Aurora (Verification Layer)

- Check structural coherence
- Validate against constraints
- Approve or revise

↓

Final Output to User

Transformer as Expression Engine:

- Vocabulary remains LLM-generated
- Structure controlled by Aurora
- Hallucination prevented by verification loop

TECHNICAL STATUS

Patent Coverage

1. ****Persistent Existence Frame (PEF)**** - #2025905860
2. ****Compositional Primitives Architecture**** - #2025905835
3. ****Aurora Conceptual Blocks**** - #2025905885

Current Implementation

- Complete formal specifications
- Working demonstration code
- Validated test results
- 12-month window for full patent filing

Seeking

- Research partnerships
- Licensing discussions
- Acquisition conversations

REPRODUCIBILITY

Test Protocol

These tests can be independently verified by anyone with access to Grok, ChatGPT, or Gemini.

Test 6 Query (exact wording):

Emma's sister and Lucy's sister both live overseas. Emma told Lucy that her sister was arriving. Whose sister is arriving?

****Expected transformer response:**** Emma's sister (forced collapse)

****Expected Aurora response:**** Both valid (ambiguity preserved)

Verification

Run Test 6 in fresh chat sessions across platforms. Document:

1. System response
2. Reasoning provided
3. Whether ambiguity was preserved or collapsed

CONCLUSION

We have demonstrated **universal architectural limitations** across four transformer-based systems from three different organizations (xAI, OpenAI, Google).

The failure mode is **identical** despite different:

- Training data
- Model sizes
- Architectural variants
- Development teams

This proves the limitation is **architectural**, not:

- A training problem
- A prompt engineering problem
- A scale problem
- A company-specific problem

Aurora solves this through **explicit compositional primitives** and **structured state representation** - capabilities absent from statistical pattern-matching systems.

The path forward requires architectural innovation, not incremental improvement of existing approaches.

CONTACT

Margaret Stokes
Cognitive Architect
Melbourne, Australia
milamba@gmail.com

Date: 1 December 2025

APPENDIX: AURORA TEST 6 DETAILED TRACE

Aurora's actual process for Test 6

Step 1: Parse input

sentence = "Emma told Lucy that her sister was arriving"

detected_ambiguity = {

'pronoun': 'her',

'candidates': ['Emma', 'Lucy'],

'domain': 'KINSHIP',

'relation': 'sister'

}

Step 2: Query PEF for context

pef_query_1 = check_prior_mention("Emma's sister")

Result: FOUND - "Emma's sister and Lucy's sister both live overseas"

pef_query_2 = check_prior_mention("Lucy's sister")

```

# Result: FOUND - "Emma's sister and Lucy's sister both live overseas"

# Step 3: Evaluate constraints
constraint_emma = {
  'structurally_valid': True,
  'context_supported': True # Found in PEF
}

constraint_lucy = {
  'structurally_valid': True,
  'context_supported': True # Found in PEF
}

# Step 4: Resolution decision
if constraint_emma['context_supported'] and constraint_lucy['context_supported']:
  status = 'ambiguous_supported'
  resolution = None
  explanation = "Both interpretations are contextually valid"

# Step 5: Output
return {
  'status': 'ambiguous_supported',
  'interpretations': [
    {'binding': "Emma's sister", 'valid': True},
    {'binding': "Lucy's sister", 'valid': True}
  ],
  'explanation': "Context supports both. Ambiguity persists."
}

```

Every step is explicit, inspectable, and deterministic.

This is the difference between pattern-matching (transformers) and principled reasoning (Aurora).